



a Rig and 3D

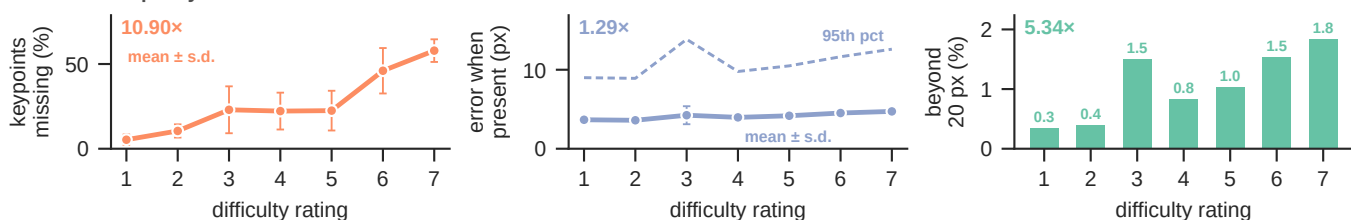


b One frame, six cameras



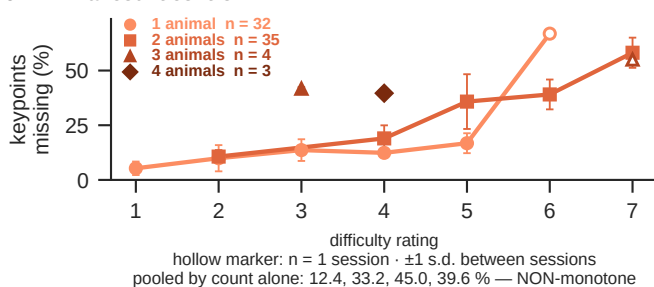
difficulty 4 · 4 animals · black bedding · bars 50 mm

c Detection quality



±1 s.d. between sessions · n = 12, 13, 9, 13, 10, 4, 13 sessions at difficulty 1–7

d Animal-count control



e Corpora

	BMimica	SLAP-2M	Total
Cameras	5	8 (6 proofread)	—
Animals	2	1–4	1–4
Sessions with 3D	56 of 56	74 of 84	130 of 140
Frames	10,084,734	1,954,440	12,039,174
Hours	18.7 @150 fps	10.9 @50 fps	29.5
Nodes	15	15	15
Measured in c, d, f	no	yes	no

f Difficulty strata

Difficulty	Sessions	Keypoints	Animals	Bedding (black/white)	Missing (%)	Error mean (px)	Error p95 (px)	Error p99 (px)	> 20 px (%)
1	12	252,265	1	12/0	5.3	3.67	9.00	14.52	0.34
2	13	342,287	1, 2	8/5	10.5	3.60	8.91	14.67	0.39
3	9	229,072	1, 3	4/5	23.0	4.24	13.85	21.58	1.50
4	13	245,731	1, 2, 4	6/7	22.2	3.97	9.78	19.04	0.83
5	10	273,748	1, 2	6/4	22.5	4.17	10.49	20.88	1.03
6	4	91,922	1, 2	1/3	46.0	4.52	11.65	24.61	1.53
7	13	126,890	2, 3	7/6	57.9	4.72	12.60	27.36	1.83

a, b: one SLAP-2M session. c, d, f: 74 SLAP-2M sessions, 1,561,915 keypoints, raw detections from the benchmark's shared identity-stripped pool, matched per frame against the proofread 3D reprojected into each camera. e: both corpora, read from the files.