


shipped — pose/cross-view-tracker.js as it ships

M1 + stale 20 + distThresh 25 — EXPERIMENTAL:

[figs/fig8-bench/xv_experimental.js](#), not in the shipped app

beside each mark: the rate, the RAW total in parentheses, then the change vs shipped. THREE DENOMINATORS, named apart, because two of them have been called “frames”: 674,111 VIDEO FRAMES × 6 cameras = 4,044,666 CAMERA-FRAMES — min(gt, det) per camera, clipped to session length — holding 7,374,061 LABELLED DETECTIONS

OTHER SLAP-2M POOL (predictions_h5s, what Fig 7 b-g use), same 42 sessions: misgrouped 303,987 -> 125,142 (-59%), switches 3,094 -> 1,312, cross-view IDF1 0.704 -> 0.721. So the mislabelled-mass result is POOL-DEPENDENT: this arm cuts switches on both pools, mass only there — never mix the pools' levels and the pre-correction claim that this arm made mass WORSE was an artefact of the broken metric, not a finding
corr3dWeight 12, Mouse-Dyad-10M's winner (371 vs 413 switches there), is NEUTRAL there: 1,314 vs 1,312 switches, IDF1 0.7205 vs 0.7212 — corpus-specific
Mouse-Dyad-10M reference (50 sessions, 5 cameras): shipped 2,071 switches (0.00460% of camera-frames) -> stale 20's 413 (0.00092%); cross-view IDF1 0.749 -> 0.861



READ THE MULTI-ANIMAL COHORT: 32 of the 74 sessions hold ONE animal and contribute 0 switches and 0 misgrouped detections, so they are EXCLUDED — pooling changed only the denominator (43%).
misgrouped = a LABELLED DETECTION — one animal, one camera, one frame — over a GT box (IoU ≥ 0.5) whose id disagrees with that box under the OPTIMAL tracker-id -> GT-index permutation for that camera, solved with linear_sum_assignment exactly as IDF1 does internally — 11.5% -> 11.7% of the labelled detections here, a LEVEL readable against IDF1

